

# Introduction to Distributed/Federated Machine Learning

---

Li Ju

PhD student@TDB

li.ju@it.uu.se

# The Optimization Problem $\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$

How to solve: **first order methods**, e.g. gradient descent:

Iteratively, at step  $t$ :  $\theta_t = \theta_t - \eta \nabla_{\theta} \sum_{n=1}^N \ell(x_n; \theta_{t-1})$

\* In practice, we use stochastic gradient descent (sgd):


$$\nabla_{\theta} \sum_{n=1}^N \ell(x_n; \theta) \approx \nabla_{\theta} \sum_{x \in \mathcal{B}} \ell(x; \theta)$$










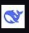

where  $\mathcal{B} \sim \{x_n\}_{n=1}^N$

# Why Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

Bigggggger models and datasets bring  
better performance (Scaling Law)



Organization ↕	Model ↕	License ↕	Parameters (B) ↕
	o3	Proprietary	-
	Claude 3.7 Sonnet	Proprietary	-
	Grok-3	Proprietary	-
	Grok-3 Mini	Proprietary	-
	o3-mini	Proprietary	-
	o1-pro	Proprietary	-
	o1	Proprietary	-
	Gemini 2.0 Flash Thinking	Proprietary	-
	o1-preview	Proprietary	-
	DeepSeek-R1	Open ©	671
	GPT-4.5	Proprietary	-

# Why Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

## Physical infeasibility












GPU RAM required:

~6 TB (batch size: 512)

Best GPU:

Nvidia H100, with 80GB

We need ~ 80 H100


Organization ↕	Model ↕	License ↕	Parameters (B) ↕
	o3	Proprietary	-
	Claude 3.7 Sonnet	Proprietary	-
	Grok-3	Proprietary	-
	Grok-3 Mini	Proprietary	-
	o3-mini	Proprietary	-
	o1-pro	Proprietary	-
	o1	Proprietary	-
	Gemini 2.0 Flash Thinking	Proprietary	-
	o1-preview	Proprietary	-
	DeepSeek-R1	Open ©	671
	GPT-4.5	Proprietary	-

# How to do Distributed ML


$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

Data parallelization: split **batch** across **M GPUs**  $\mathcal{B} = \cup_{m=1}^M \mathcal{B}_m$


Parameter  
server




Allreduce



Ring Allreduce



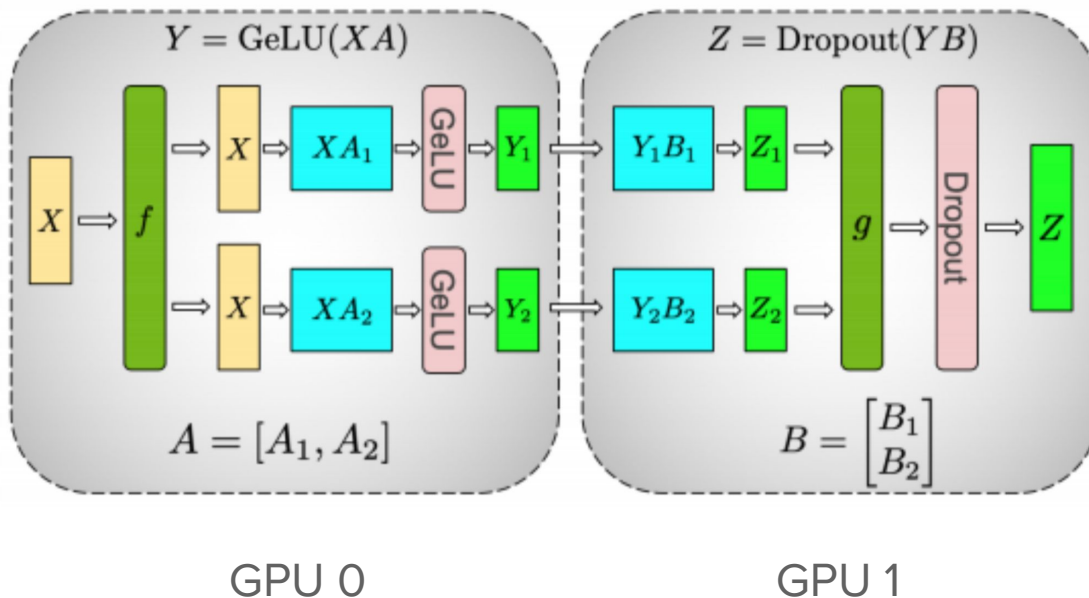
Tree Allreduce



# How to do Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

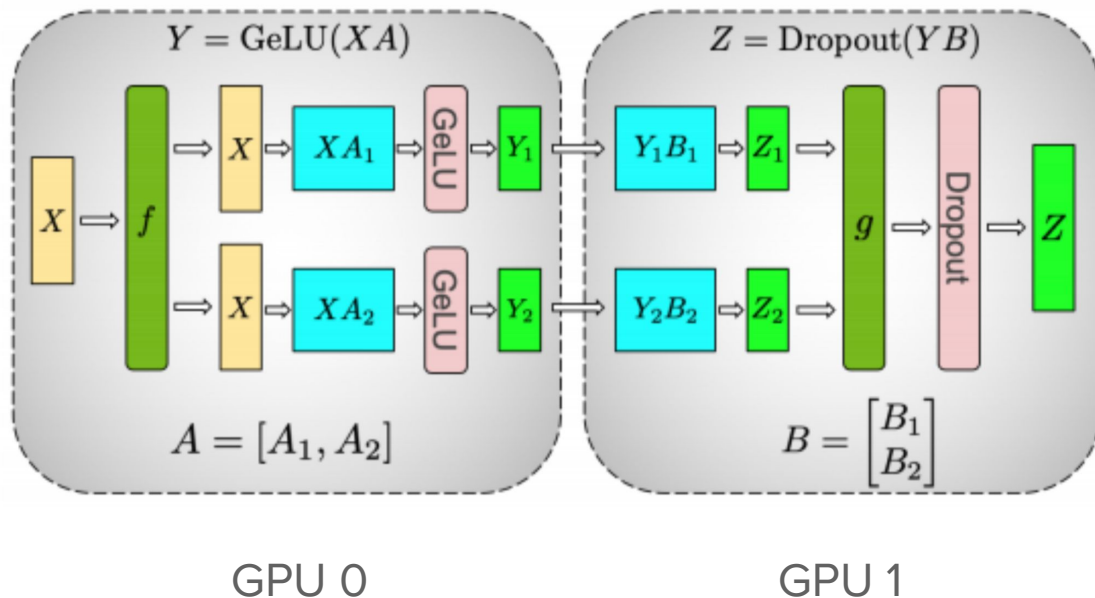
Model parallelization?



# How to do Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

## Model parallelization?



This is **NOT** model parallelization!


Workloads on GPUs depend on each other...

It is called **workload partitioning**.

# How to do Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

Pipeline parallelization:



\* The batch of data needs to be split into mini batches.



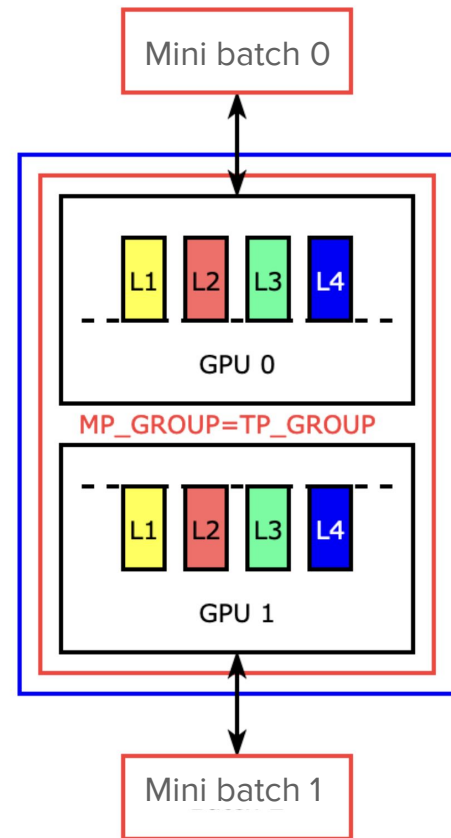

# How to do Distributed ML

## Model Parallelization:

- Zero Redundancy Optimizer (ZeRO)
- Tensor parallelization
- ...

ZeRO is the most common approach.

\* The batch of data needs to be split into M mini batches.




# How to do Distributed ML

$$\arg_{\theta} \min \sum_{n=1}^N \ell(x_n; \theta)$$

## 3D-parallelization

Data, pipeline and model

parallelization are orthogonal to  
each other:



# Backend

Taking PyTorch as an example:

MPI: CPU


NCCL: Nvidia GPU

GLOO: Both (partially)


Backend	gloo		mpi		nccl	
Device	CPU	GPU	CPU	GPU	CPU	GPU
send	✓	✗	✓	?	✗	✓
recv	✓	✗	✓	?	✗	✓
broadcast	✓	✓	✓	?	✗	✓
all_reduce	✓	✓	✓	?	✗	✓
reduce	✓	✗	✓	?	✗	✓
all_gather	✓	✗	✓	?	✗	✓
gather	✓	✗	✓	?	✗	✓

# Why Federated ML

## Intrinsic distributed nature + privacy concern



# How to do Federated ML



I. Initialize a seed model

II. Iteratively:

1. Broadcast model params
2. Training locally
3. Collecting local models
4. Aggregating and updating model

# How to do Federated ML

---

**Algorithm 1** Federated Averaging (FEDAVG)

---

```
1: procedure FEDAVG ( $\mathbf{x}^{(0,0)}, \eta$ )
2: for  $r = 0, \dots, R - 1$  do
3:   on client for  $m \in [M]$  in parallel do
4:      $\mathbf{x}_m^{(r,0)} \leftarrow \mathbf{x}^{(r,0)}$   $\triangleright$  broadcast current iterate
5:     for  $k = 0, \dots, K - 1$  do
6:        $\xi_m^{(r,k)} \sim \mathcal{D}_m$ 
7:        $\mathbf{g}_m^{(r,k)} \leftarrow \nabla f(\mathbf{x}_m^{(r,k)}; \xi_m^{(r,k)})$ 
8:        $\mathbf{x}_m^{(r,k+1)} \leftarrow \mathbf{x}_m^{(r,k)} - \eta \cdot \mathbf{g}_m^{(r,k)}$   $\triangleright$  client update
9:    $\mathbf{x}^{(r+1,0)} \leftarrow \frac{1}{M} \sum_{m=1}^M \mathbf{x}_m^{(r,K)}$   $\triangleright$  server averaging
```

---


$K=1$ : Parameter-server data parallelization /  
Distributed SGD

$K \gg 1$ : Federated learning

# Why FML is significant

**Intrinsic mathematical issues:** Distributed SGD is equivalent to SGD, but FedAvg is not:


- Different local steps
- Size heterogeneity of local data
- Statistical heterogeneity of local data
- Convergence analysis
- ...



# Why FML is significant

## Practical/engineering issues:

- Extremely bottlenecked by communication
- Straggler problem
- Asynchronized FML
- Personalization
- Machine Unlearning (quitting participant)
- ...





Any Questions?